

1. Advantages and Performance of the Relational Model

- **Standardization and Familiarity:** The relational model provides a "(Mostly) Standard Data Model and Query Language," ensuring broad comprehension and support.
- **ACID Compliance:** A key strength of relational databases is their adherence to ACID properties, which guarantee reliable transaction processing:
 - **Atomicity:** Transactions are treated as single, indivisible units—either fully executed or not executed at all.
 - **Consistency:** Transactions transition the database from one valid state to another, ensuring all integrity constraints are met.
 - **Isolation:** Simultaneously executing transactions do not interfere with one another. Potential issues such as "Dirty Read," "Non-repeatable Read," and "Phantom Reads" are discussed.
 - **Durability:** Once committed, transaction changes are permanently recorded.
- **Optimized Data Handling:** RDBMSs enhance efficiency through techniques like:
 - Indexing
 - Storage control (row vs. column-oriented)
 - Query optimization
 - Caching and prefetching
 - Materialized views
 - Precompiled stored procedures
 - Data replication and partitioning

2. Transaction Processing and ACID Properties in Depth

- **Defining a Transaction:** Transactions consist of one or more CRUD operations executed as a single logical unit of work. They either complete successfully (**COMMIT**) or fully revert changes (**ROLLBACK** or **ABORT**) in case of failure.
- **Significance of Transactions:** They are essential for:
 - Maintaining Data Integrity
 - Error Recovery
 - Managing Concurrency
 - Ensuring Reliable Data Storage
 - Simplifying Error Handling
- **Illustrative Example:** SQL code for a **transfer** procedure showcases transaction implementation using **START TRANSACTION**, **UPDATE**, **INSERT**, **ROLLBACK** (e.g., with an error message: "Transaction rolled back: Insufficient funds"), and **COMMIT**, demonstrating atomicity and consistency.

3. Limitations of the Relational Model in Contemporary Applications

- **Challenges in Schema Evolution:** Frequent structural changes can be difficult to manage.

- **Overhead for Certain Applications:** Some use cases do not require strict ACID compliance.
- **Expensive Join Operations:** Complex joins across multiple tables can degrade performance.
- **Handling Non-Relational Data:** Relational databases are not naturally suited for unstructured and semi-structured formats like JSON and XML.
- **Scaling Constraints:** Horizontal scaling (distributing databases across multiple servers) is traditionally complex.
- **Performance Demands:** Certain applications require lower latency and higher throughput than conventional RDBMSs can consistently provide.

4. Transition to Distributed Systems and Data Storage

- **Scaling Approaches:** Initially, vertical scaling (enhancing hardware) was preferred due to simplicity. However, cost and availability constraints necessitate horizontal scaling (distributing workloads across multiple machines).
- **Definition of a Distributed System:** As per Andrew Tannenbaum, a distributed system is "a collection of independent computers that appear to its users as one computer."
- **Characteristics of Distributed Systems:**
 - Concurrent operation across multiple computers
 - Independent system failures
 - Lack of a shared global clock
- **Distributed Storage Strategies:**
 - **Single Main Node:** Traditional centralized storage.
 - **Distributed Data Stores:** Data is replicated across multiple nodes to enhance redundancy.
 - **Distributed Databases:** Can be relational (e.g., MySQL, PostgreSQL with replication/sharding, CockroachDB) or non-relational (NoSQL solutions).
- **Network Partitioning:** Failures in network connectivity are inevitable, necessitating systems that remain operational under partitioned conditions.

5. The CAP Theorem: Trade-offs in Distributed Data Storage

- **Definition:** The CAP Theorem asserts that a distributed data store cannot simultaneously provide all three of the following guarantees:
 - **Consistency:** Ensures every read operation retrieves the most recent write or returns an error. (Distinct from ACID consistency.)
 - **Availability:** Guarantees that every request receives a response, though it may not reflect the latest data.
 - **Partition Tolerance:** Ensures the system remains operational despite network failures.
- **Database Perspective on CAP:**
 - **Consistency:** All users perceive an identical dataset at any given moment.
 - **Availability:** The database remains operational even when failures occur.

- **Partition Tolerance:** The database continues functioning despite network partitioning.
- **Trade-offs:**
 - **Consistency + Availability:** Vulnerable to network partitions.
 - **Consistency + Partition Tolerance:** May reduce availability by rejecting requests during partitions.
 - **Availability + Partition Tolerance:** Can lead to outdated data being served due to lack of consistency.
- **Practical Implications:** While often simplified as a necessity to sacrifice one property, the reality is more nuanced. In the presence of network failures, achieving both consistency and availability simultaneously becomes impossible.